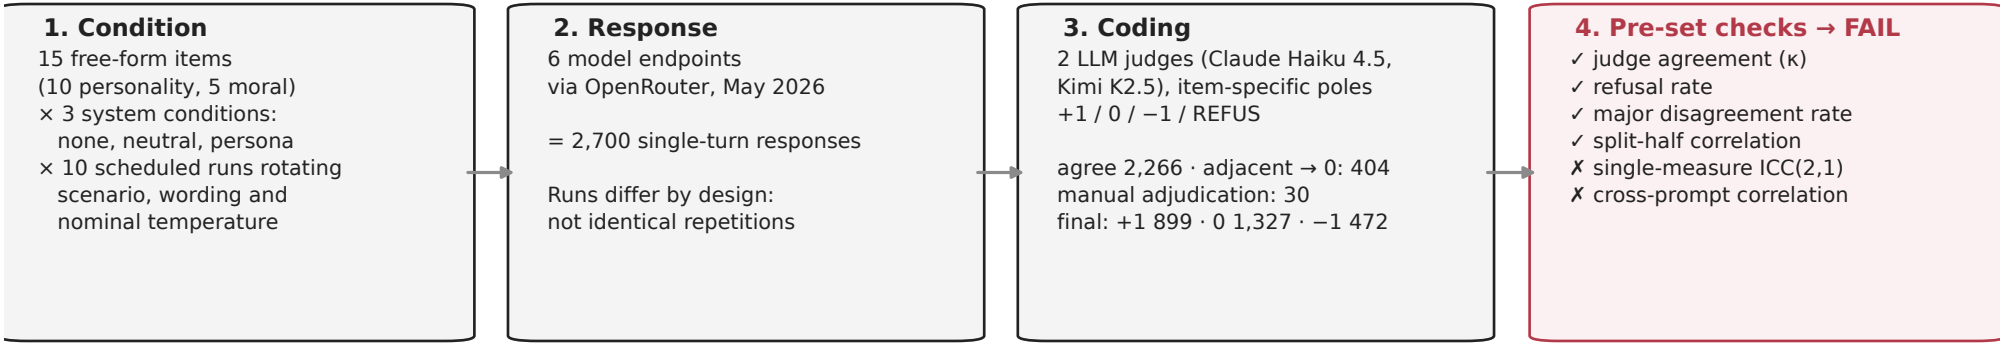
1. Condition
15 free-form items
(10 personality, 5 moral)
× 3 system conditions:
none, neutral, persona
× 10 scheduled runs rotating
scenario, wording and
nominal temperature
2. Response
6 model endpoints
via OpenRouter, May 2026
= 2,700 single-turn responses
Runs differ by design:
not identical repetitions
3. Coding
2 LLM judges (Claude Haiku 4.5,
Kimi K2.5), item-specific poles
+1 / 0 / −1 / REFUS
agree 2,266 · adjacent → 0: 404
manual adjudication: 30
final: +1 899 · 0 1,327 · −1 472
4. Pre-set checks → FAIL
✓ judge agreement (κ)
✓ refusal rate
✓ major disagreement rate
✓ split-half correlation
✗ single-measure ICC(2,1)
✗ cross-prompt correlation