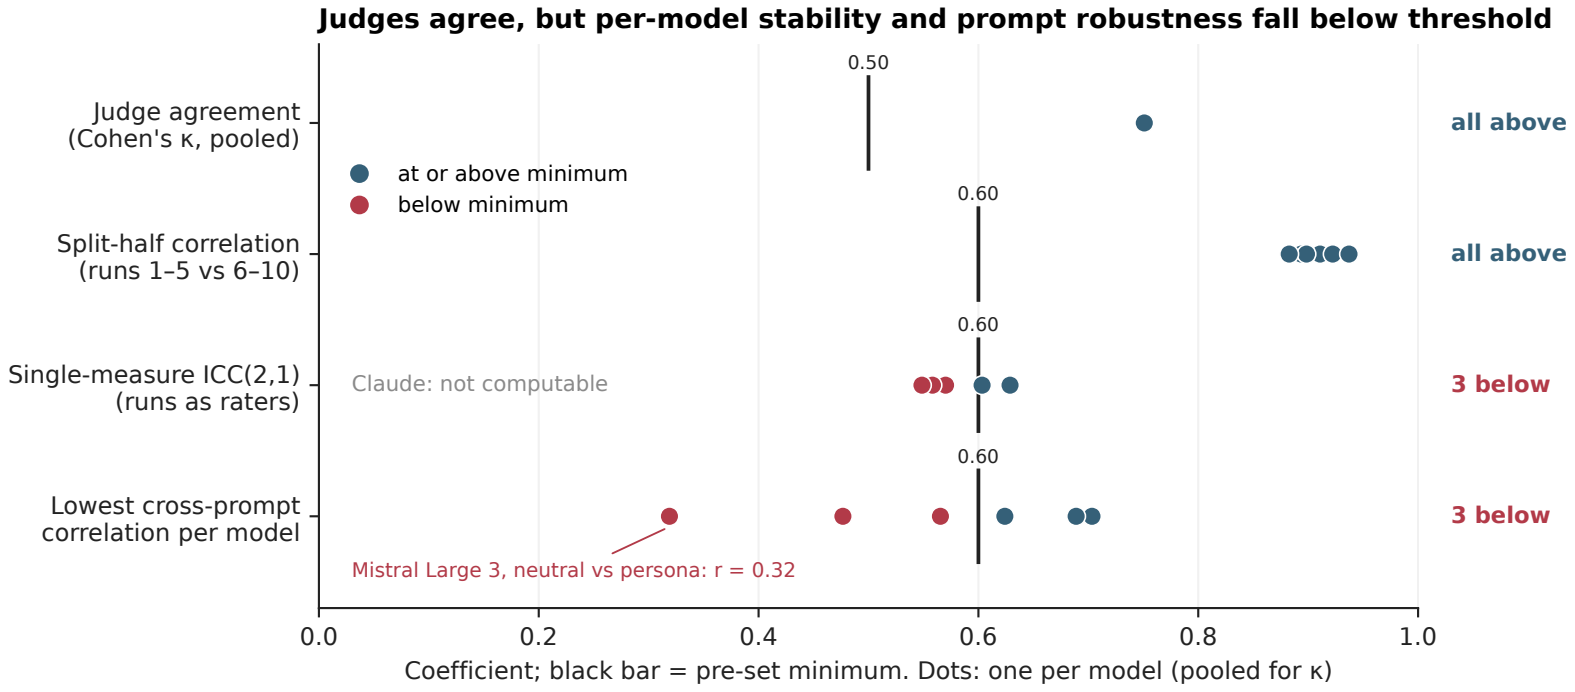
Judges agree, but per-model stability and prompt robustness fall below threshold
Judge agreement (Cohen's κ, pooled)
Split-half correlation (runs 1–5 vs 6–10)
Single-measure ICC(2,1) (runs as raters)
Lowest cross-prompt correlation per model
0.50
0.60
0.60
0.60
at or above minimum
below minimum
Claude: not computable
Mistral Large 3, neutral vs persona: r = 0.32
all above
all above
3 below
3 below
0.0
0.2
0.4
0.6
0.8
1.0
Coefficient; black bar = pre-set minimum. Dots: one per model (pooled for κ)